

## ✓ DSBDAL Assignment 09 - Data Analytics 2




1. Implement logistic regression using Python/R to perform classification on Social\_Network\_Ads.csv dataset.
2. Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall on the given dataset

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.r

```
import numpy as np
import pandas as pd
import seaborn as sns
```

```
ds = pd.read_csv('/content/drive/My Drive/DSBDL/Assignment9/data.csv')
ds
```

	User ID	Gender	Age	EstimatedSalary	Purchased	
0	15624510	Male	19	19000	0	
1	15810944	Male	35	20000	0	
2	15668575	Female	26	43000	0	
3	15603246	Female	27	57000	0	
4	15804002	Male	19	76000	0	
...	...	...	...	...	...	
395	15691863	Female	46	41000	1	
396	15706071	Male	51	23000	1	
397	15654296	Female	50	20000	1	
398	15755018	Male	36	33000	0	
399	15594041	Female	49	36000	1	

400 rows × 5 columns

Next steps:

[Generate code with ds](#)[View recommended plots](#)



```
ds.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID               400 non-null   int64
1   Gender                400 non-null   object
2   Age                   400 non-null   int64
3   EstimatedSalary       400 non-null   int64
4   Purchased             400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB

```

```
ds.describe()
```

	User ID	Age	EstimatedSalary	Purchased	
<b>count</b>	4.000000e+02	400.000000	400.000000	400.000000	
<b>mean</b>	1.569154e+07	37.655000	69742.500000	0.357500	
<b>std</b>	7.165832e+04	10.482877	34096.960282	0.479864	
<b>min</b>	1.556669e+07	18.000000	15000.000000	0.000000	
<b>25%</b>	1.562676e+07	29.750000	43000.000000	0.000000	
<b>50%</b>	1.569434e+07	37.000000	70000.000000	0.000000	
<b>75%</b>	1.575036e+07	46.000000	88000.000000	1.000000	
<b>max</b>	1.581524e+07	60.000000	150000.000000	1.000000	

```
ds.isna().sum()
```

```

User ID      0
Gender       0
Age          0
EstimatedSalary  0
Purchased    0
dtype: int64




```

```
from sklearn.preprocessing import LabelEncoder
```

```

ds['Gender'] = LabelEncoder().fit_transform(ds['Gender'])
ds

```

	User ID	Gender	Age	EstimatedSalary	Purchased	
0	15624510	1	19	19000	0	
1	15810944	1	35	20000	0	
2	15668575	0	26	43000	0	
3	15603246	0	27	57000	0	
4	15804002	1	19	76000	0	
...	...	...	...	...	...	
395	15691863	0	46	41000	1	
396	15706071	1	51	23000	1	
397	15654296	0	50	20000	1	
398	15755018	1	36	33000	0	
399	15594041	0	49	36000	1	



400 rows × 5 columns

Next steps:

[Generate code with ds](#)

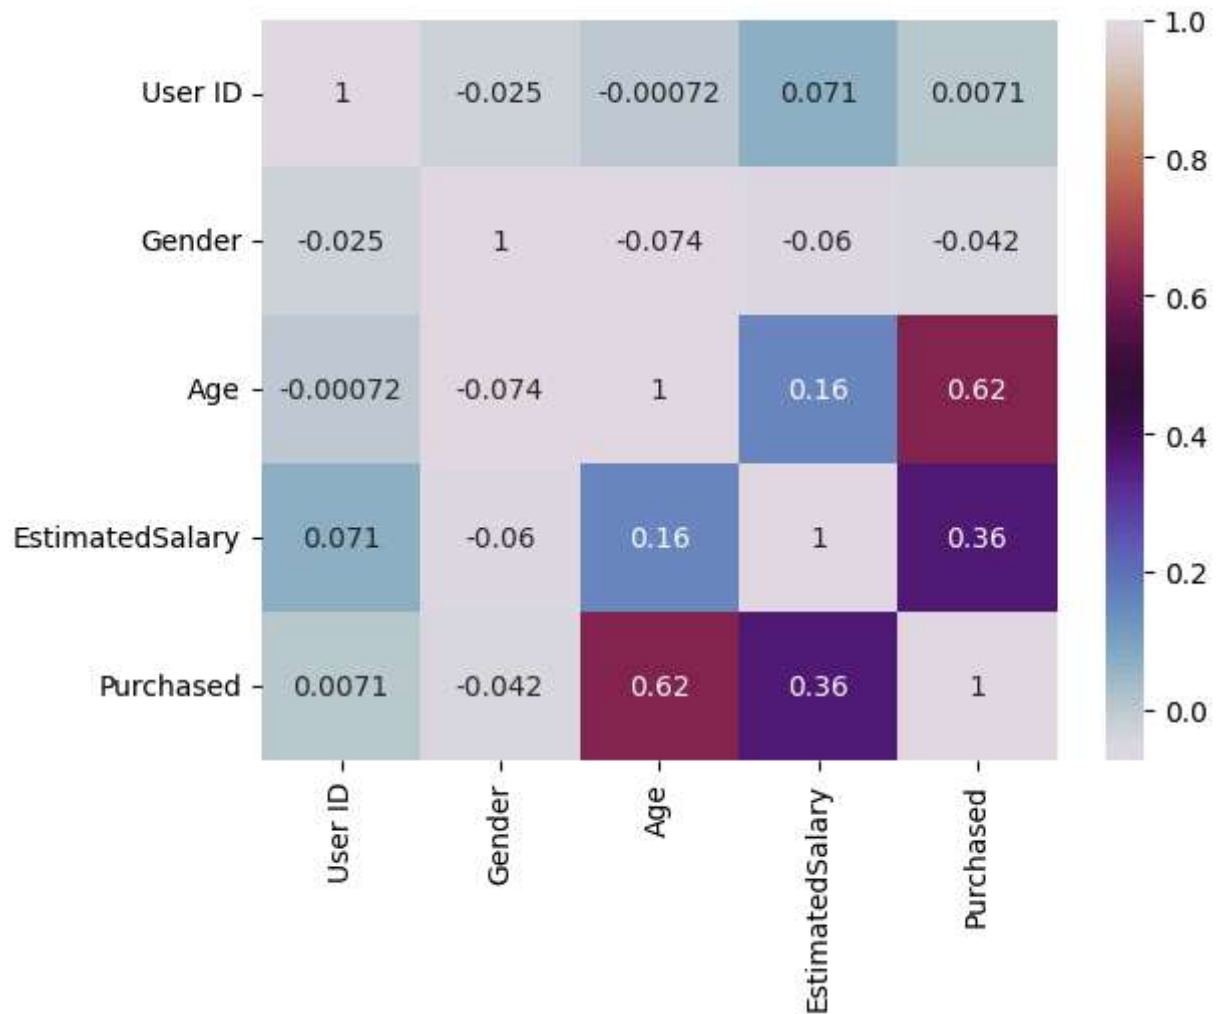
 [View recommended plots](#)

`ds.corr()`

	User ID	Gender	Age	EstimatedSalary	Purchased	
<b>User ID</b>	1.000000	-0.025249	-0.000721	0.071097	0.007120	
<b>Gender</b>	-0.025249	1.000000	-0.073741	-0.060435	-0.042469	
<b>Age</b>	-0.000721	-0.073741	1.000000	0.155238	0.622454	
<b>EstimatedSalary</b>	0.071097	-0.060435	0.155238	1.000000	0.362083	
<b>Purchased</b>	0.007120	-0.042469	0.622454	0.362083	1.000000	

`sns.heatmap(ds.corr(), cmap = 'twilight', annot = True)`

<Axes: >



```
from sklearn.model_selection import train_test_split
```

```
ds.drop('User ID', axis = 1, inplace = True)
```

```
X = np.asarray(ds.drop('Purchased', axis = 1))
```

```
y = np.asarray(ds['Purchased'])
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state
```

```
from sklearn.preprocessing import StandardScaler
```

```
X_train = StandardScaler().fit_transform(X_train)
```

```
X_test = StandardScaler().fit_transform(X_test)
```

```
from sklearn.linear_model import LogisticRegression
```

```
regressor = LogisticRegression()
```

```
regressor.fit(X_train, y_train)
```

```
y_pred = regressor.predict(X_test)
```

```
print(y_pred)
```

```
[0 1 0 1 0 0 1 0 0 0 0 1 0 0 0 0 1 1 0 1 0 0 0 1 0 0 1 0 1 0 0 0 1 0 1 0 0  
 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0]
```

```
0 0 1 1 0 0 0 0 1 0 1 0 0 1 0 0 1 0 0 0 0 0 1 0 1 0]
```

```
from sklearn.metrics import confusion_matrix, classification_report
```

```
matrix = confusion_matrix(y_test, y_pred)
```

```
print(matrix)
```

```
print(classification_report(y_test, y_pred))
```

```
[[61  2]
 [12 25]]
```

	precision	recall	f1-score	support
0	0.84	0.97	0.90	63
1	0.93	0.68	0.78	37
accuracy			0.86	100
macro avg	0.88	0.82	0.84	100
weighted avg	0.87	0.86	0.85	100

```
from sklearn.metrics import confusion_matrix, accuracy_score, f1_score, precision_score, re
```

```
matrix = confusion_matrix(y_test, y_pred)
```

```
tn, fp, fn, tp = matrix.ravel()
```

```
print("True Positives: ", tp)
```

```
print("True Negatives: ", tn)
```

```
print("False Positives: ", fp)
```

```
print("False Negatives: ", fn)
```

```
print()
```

```
print("Accuracy: ", (tp + tn) / (tp + tn + fn + fp))
```

```
print("Recall: ", tp / (tp + fn))
```

```
print("Precision: ", tp / (tp + fp))
```

```
True Positives: 25
True Negatives: 61
False Positives: 2
False Negatives: 12
```

```
Accuracy: 0.86
Recall: 0.6756756756756757
Precision: 0.9259259259259259
```

```
print("Accuracy: ", accuracy_score(y_test, y_pred))
```

```
print("Recall: ", recall_score(y_test, y_pred))
```

```
print("Precision: ", precision_score(y_test, y_pred))
```



```
Accuracy: 0.86
Recall: 0.6756756756756757
Precision: 0.9259259259259259
```